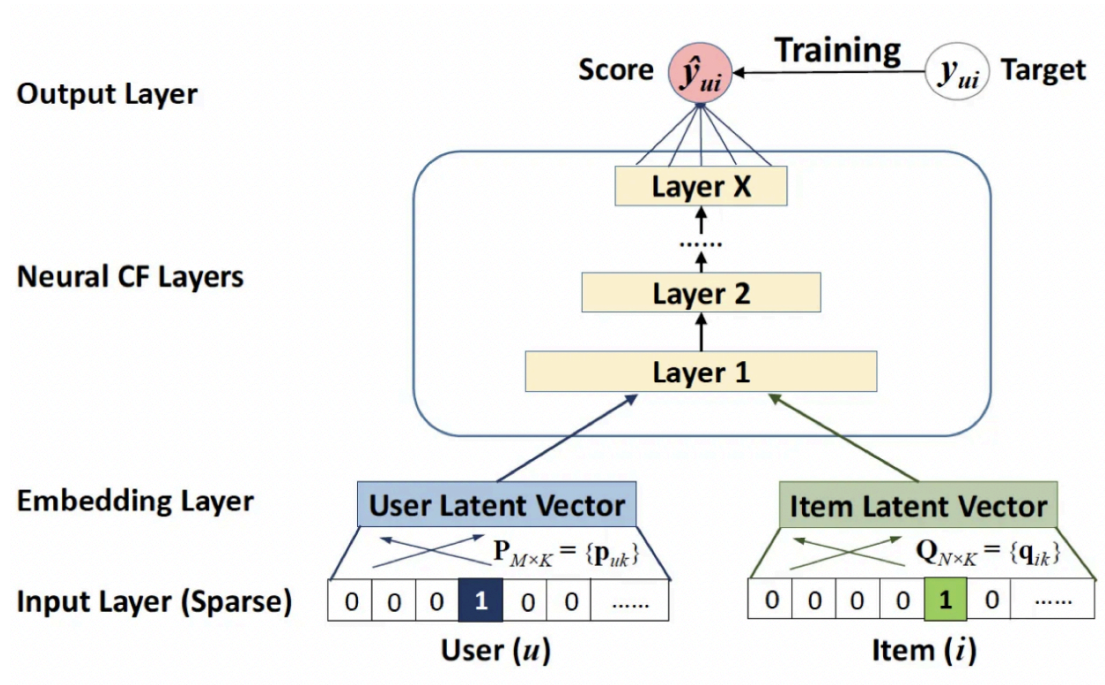




Neural Collaborative Filtering

NCF Architecture



Data Loading And Preprocessing

```
In [1]: import pandas as pd
import numpy as np

data = pd.read_csv("/kaggle/input/movielens-32m/ml-32m/ratings.csv")
data.drop(columns=["timestamp"], inplace=True)
data.head()
```

```
Out[1]:
```

	userId	movieId	rating
0	1	17	4.0
1	1	25	1.0
2	1	29	2.0
3	1	30	5.0
4	1	32	5.0

```
In [2]: user_mapping = {user_id: idx for idx, user_id in enumerate(data['userId'].unique())}
item_mapping = {item_id: idx for idx, item_id in enumerate(data['movieId'].unique())}

data['userId'] = data['userId'].map(user_mapping)
data['movieId'] = data['movieId'].map(item_mapping)

num_users = len(user_mapping)
```

```
num_items = len(item_mapping)

print(f"Numbers of users: {num_users}\nNumber of movies: {num_items}")
```

Numbers of users: 200948
Number of movies: 84432

```
In [3]: from sklearn.model_selection import train_test_split

x = data[['userId', 'movieId']].values
y = (data['rating'] > 3).astype(np.float32)

X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_

print(f"Training samples: {len(X_train)}, Testing samples: {len(X_test)}")
```

Training samples: 25600163, Testing samples: 6400041

```
In [4]: import tensorflow as tf

train_dataset = tf.data.Dataset.from_tensor_slices((X_train, y_train))
test_dataset = tf.data.Dataset.from_tensor_slices((X_test, y_test))

batch_size = 256

train_dataset = train_dataset.shuffle(buffer_size=len(X_train)).batch(batch_size)
test_dataset = test_dataset.batch(batch_size)
```

```
In [5]: def preprocess(data):
    user_ids = data[:, 0]
    item_ids = data[:, 1]
    return {"user_id": user_ids, "item_id": item_ids}

train_dataset = train_dataset.map(lambda x, y: (preprocess(x), y))
test_dataset = test_dataset.map(lambda x, y: (preprocess(x), y))
```

```
In [6]: for batch in train_dataset.take(1):
    inputs, labels = batch
    print("User IDs:", inputs['user_id'][:5].numpy())
    print("Item IDs:", inputs['item_id'][:5].numpy())
    print("Labels:", labels[:5].numpy())
```

User IDs: [74358 52531 40164 197739 14798]
Item IDs: [123 514 785 1295 247]
Labels: [1. 1. 1. 1. 1.]

Model Creation

```
In [7]: from tensorflow.keras import layers, Model
from tensorflow.keras.saving import register_keras_serializable

@register_keras_serializable(package="NCFModel")
class NCFModel(Model):
    def __init__(self, num_users, num_items, embedding_dim=16, mlp_layers=[64, 3
        super(NCFModel, self).__init__(**kwargs)

        self.num_users = num_users
        self.num_items = num_items
        self.embedding_dim = embedding_dim
```

```

self.mlp_layers_config = mlp_layers

self.user_embedding_gmf = layers.Embedding(input_dim = num_users, output
self.item_embedding_gmf = layers.Embedding(input_dim = num_items, output

self.user_embedding_mlp = layers.Embedding(input_dim = num_users, output
self.item_embedding_mlp = layers.Embedding(input_dim = num_items, output

self.mlp_layers = tf.keras.Sequential()
for units in mlp_layers:
    self.mlp_layers.add(layers.Dense(units, activation='relu'))

self.concat_layer = layers.Concatenate()
self.output_layer = layers.Dense(1, activation='sigmoid')

def call(self, inputs):
    user_ids, item_ids = inputs["user_id"], inputs['item_id']
    gmf_user = self.user_embedding_gmf(user_ids)
    gmf_item = self.item_embedding_gmf(item_ids)
    gmf_output = gmf_user * gmf_item

    mlp_user = self.user_embedding_mlp(user_ids)
    mlp_item = self.item_embedding_mlp(item_ids)
    mlp_input = tf.concat([mlp_user, mlp_item], axis=-1)
    mlp_output = self.mlp_layers(mlp_input)

    combined = self.concat_layer([gmf_output, mlp_output])

    return self.output_layer(combined)
def get_config(self):
    # Returns the configuration to save the model
    config = super().get_config()
    config.update({
        "num_users": self.num_users,
        "num_items": self.num_items,
        "embedding_dim": self.embedding_dim,
        "mlp_layers": self.mlp_layers_config,
    })
    return config

@classmethod
def from_config(cls, config):
    # Creates the model from the saved configuration
    return cls(
        num_users=config["num_users"],
        num_items=config["num_items"],
        embedding_dim=config["embedding_dim"],
        mlp_layers=config["mlp_layers"],
    )

```

```
In [8]: ncf_model = NCFModel(num_users=200948, num_items=84432, embedding_dim=16, mlp_la
```

```
In [9]: ncf_model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
    loss='binary_crossentropy',
    metrics=['accuracy', 'mse', 'mae']
)
```

```
In [10]: ncf_model.summary()
```

Model: "ncf_model"

Layer (type)	Output Shape	Param #
embedding (Embedding)	?	0 (unbuilt)
embedding_1 (Embedding)	?	0 (unbuilt)
embedding_2 (Embedding)	?	0 (unbuilt)
embedding_3 (Embedding)	?	0 (unbuilt)
sequential (Sequential)	?	0 (unbuilt)
concatenate (Concatenate)	?	0 (unbuilt)
dense_3 (Dense)	?	0 (unbuilt)



Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)

Non-trainable params: 0 (0.00 B)

Training The Model

```
In [11]: ncf_model.fit(
          train_dataset,
          validation_data=test_dataset,
          epochs=1,
          batch_size=256
        )
```

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1732784999.443291 310 service.cc:145] XLA service 0x78f0b4058920 initialized for platform CUDA (this does not guarantee that XLA will be used). Devices:

I0000 00:00:1732784999.443349 310 service.cc:153] StreamExecutor device (0): Tesla T4, Compute Capability 7.5

I0000 00:00:1732784999.443353 310 service.cc:153] StreamExecutor device (1): Tesla T4, Compute Capability 7.5

34/100001 ————— 7:52 5ms/step - accuracy: 0.5875 - loss: 0.6864 - mae: 0.4960 - mse: 0.2466

I0000 00:00:1732785000.813665 310 device_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

100001/100001 ————— 502s 4ms/step - accuracy: 0.7432 - loss: 0.5176 - mae: 0.3451 - mse: 0.1726 - val_accuracy: 0.7740 - val_loss: 0.4699 - val_mae: 0.3090 - val_mse: 0.1542

```
Out[11]: <keras.src.callbacks.history.History at 0x78f117118c10>
```

Evaluation

```
In [13]: loss, accuracy, _, _ = ncf_model.evaluate(test_dataset)
          print(f"Loss = {loss:.4f}, Accuracy = {accuracy:.4f}")
```

25001/25001 ————— 40s 2ms/step - accuracy: 0.7741 - loss: 0.4699 -
 mae: 0.3090 - mse: 0.1542
 Loss = 0.4699, Accuracy = 0.7740

In [14]: `ncf_model.summary()`

Model: "ncf_model"

Layer (type)	Output Shape	Param #
embedding (Embedding)	?	3,215,168
embedding_1 (Embedding)	?	1,350,912
embedding_2 (Embedding)	?	3,215,168
embedding_3 (Embedding)	?	1,350,912
sequential (Sequential)	?	4,720
concatenate (Concatenate)	?	0
dense_3 (Dense)	?	33



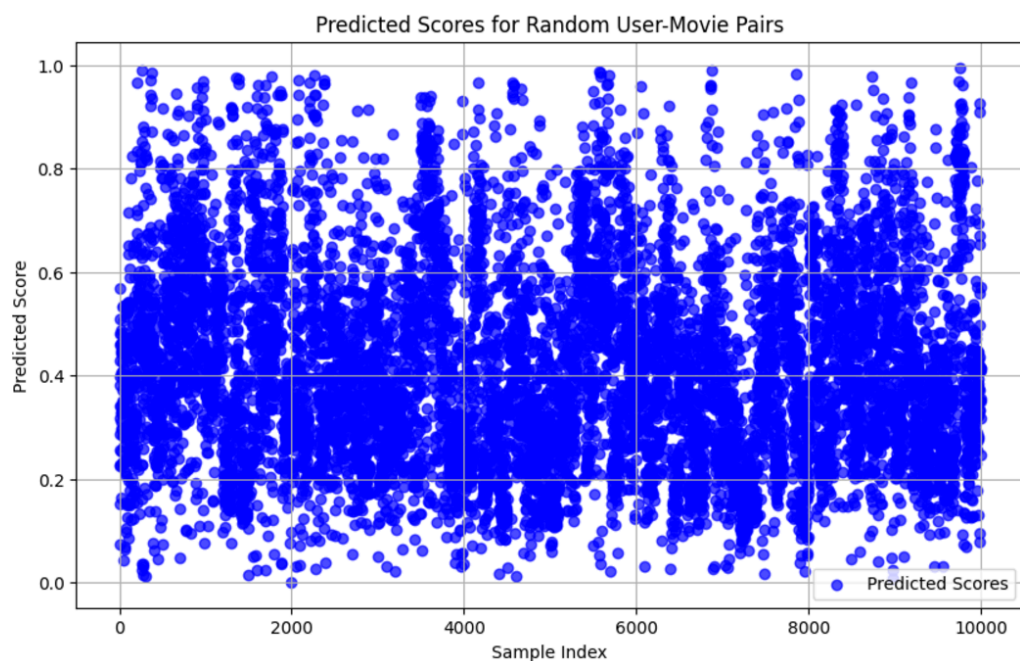
Total params: 27,410,741 (104.56 MB)

Trainable params: 9,136,913 (34.85 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 18,273,828 (69.71 MB)

In [15]: `ncf_model.save("/kaggle/working/ncf_model.keras")`



In []: